

A Learning Automata-Based Multiobjective Hyper-Heuristic

Wenwen Li^{ID}, Ender Özcan, *Senior Member, IEEE*, and Robert John, *Senior Member, IEEE*

Abstract—Metaheuristics, being tailored to each particular domain by experts, have been successfully applied to many computationally hard optimization problems. However, once implemented, their application to a new problem domain or a slight change in the problem description would often require additional expert intervention. There is a growing number of studies on reusable cross-domain search methodologies, such as selection hyper-heuristics, which are applicable to problem instances from various domains, requiring minimal expert intervention or even none. This paper introduces a new **learning automata-based selection hyper-heuristic** controlling a set of **multiobjective metaheuristics**. The approach operates above three well-known multiobjective evolutionary algorithms and mixes them, exploiting the strengths of each algorithm. The performance and behavior of **two variants of the proposed selection hyper-heuristic**, each utilizing a **different initialization scheme** are investigated across a range of unconstrained multiobjective mathematical benchmark functions from two different sets and the real-world problem of vehicle crash-worthiness. The empirical results illustrate the effectiveness of our approach for cross-domain search, regardless of the initialization scheme, on those problems when compared to each individual multiobjective algorithm. Moreover, both variants perform significantly better than some previously proposed selection hyper-heuristics for multiobjective optimization, thus significantly enhancing the opportunities for improved multiobjective optimization.

Index Terms—Evolutionary algorithms, hyper-heuristics, multiobjective optimization, online learning, operational research.

I. INTRODUCTION

MULTIOBJECTIVE optimization problems (MOPs) require simultaneous handling of various and often conflicting objectives during the search process. The solution methods designed for MOPs seek a set of “equivalent” solutions, each reflecting a *tradeoff* between different objectives.

There are distinct complexities associated with MOPs making the development of effective and efficient solution methods extremely challenging (e.g., very large search spaces, noise, uncertainty, etc.). Metaheuristics, in particular, multiobjective

evolutionary algorithms (MOEAs) are the most commonly used search methods in the area of solving MOPs. One of the main advantages of MOEAs is that they are population-based techniques, capable of obtaining a set of tradeoff solutions with reasonable quality even in a single run [1]. Even though “optimality” cannot be guaranteed, empirical results indicate the success of MOEAs on a variety of problem domains, including planning and scheduling [2], [3], data mining [4], and circuits and communications [5]. There are different types of MOEAs, each utilizing different algorithmic components during the search process and so perform differently. In the majority of the previous studies, individual MOEAs are designed and applied to a particular problem in hand. More on MOEAs and their applications to various multiobjective problems can be found in [6].

On the other hand, there is a growing number of studies on *selection hyper-heuristics* which provide a general-purpose heuristic optimization framework for utilizing the strengths of multiple (meta)heuristics [7]. Selection hyper-heuristics control and mix low level (meta)heuristics, automatically deciding which one(s) to apply to the candidate solution(s) at each decision point of the iterative search process [8]. Raising the generality level of heuristic optimization methods is one of the main motivations behind the hyper-heuristic studies. The idea is, through automation of the heuristic search, to provide effective and reusable *cross-domain* search methodologies which are applicable to the problems with different characteristics from various domains without requiring much expert involvement.

Learning is key to develop an effective selection hyper-heuristic with the adaptation capability. There are some recent studies looking into the interplay between data science techniques, particularly machine learning algorithms and selection hyper-heuristics leading to an improved overall performance. For example, Asta *et al.* [9], [10] used tensor analysis as a machine learning approach to decide which low level heuristics to employ at different stages of the search process. In [11], the feasibility and effectiveness of using reinforcement learning to improve the performance of metaheuristics and hyper-heuristics have been discussed in depth. Kheiri and Özcan [12] introduced an effective multistage hyper-heuristic for cross-domain search which first, reduces the low level heuristics to be used in the following stage based on a multiobjective learning strategy and then mixes them under a stochastic local search framework. More recently, computational intelligence techniques have been used as components of general purpose methods managing low

Manuscript received April 25, 2017; revised August 13, 2017 and October 24, 2017; accepted December 11, 2017. Date of publication December 20, 2017; date of current version January 28, 2019. (*Corresponding author: Wenwen Li.*)

The authors are with the Automated Scheduling, Optimisation and Planning Research Group, School of Computer Science, University of Nottingham, Nottingham NG8 1BB, U.K. (e-mail: psxwl8@nottingham.ac.uk; pszeo@nottingham.ac.uk; pszrij@nottingham.ac.uk).

Color versions of one or more of the figures in this paper are available online at <http://ieeexplore.ieee.org>.

Digital Object Identifier 10.1109/TEVC.2017.2785346

1089-778X © 2017 IEEE. Personal use is permitted, but republication/redistribution requires IEEE permission.
See http://www.ieee.org/publications_standards/publications/rights/index.html for more information.

level (meta)heuristics for overall performance improvement. For example, Zamli *et al.* [13] introduced a fuzzy inference selection-based hyper-heuristic which mixed and controlled four search operators, each derived from a different meta-heuristic to solve a computationally hard problem of t -way test suite generation. However, the aforementioned studies all focus on single objective optimization. There have been some studies on combining the strengths of multiple MOEAs with the aim of providing a better overall performance for multiobjective optimization under a selection hyper-heuristic framework (e.g., [14] and [15]). From this point onward, we will refer to such selection hyper-heuristics as multiobjective hyper-heuristics (MOHHs).

In this paper, we present a new learning automata-based selection hyper-heuristic framework with implementation of two variants, learning automata-based hyper-heuristic (HH-LA) and learning automata-based hyper-heuristic with a ranking scheme initialization (HH-RILA) for multiobjective optimization. Both selection hyper-heuristics mix and control a set of three well-known MOEAs: 1) nondominated sorting genetic algorithm (NSGA-II) [16]; 2) strength Pareto evolutionary algorithm 2 (SPEA2) [17]; and 3) indicator-based evolutionary algorithm (IBEA) [18]. The learning automaton acts as a guidance for choosing the appropriate MOEA at each decision point while solving a given problem. The proposed **two variants** of selection hyper-heuristics **mainly differ in their initial set-up process**. **HH-LA** employs all three low level MOEAs and gives an equal chance initially to each algorithm making a random start. **HH-RILA** applies a ranking scheme which eliminates the relatively poor performing MOEA(s) and uses the remaining MOEAs in the improvement process (Section III-A). The performance of the proposed hyper-heuristics are investigated against a variety of other multiobjective approaches across a range of multiobjective problems, including well-known benchmark functions and a real-world problem of vehicle crashworthiness. The empirical results indicate the effectiveness and generality of the proposed hyper-heuristics with novel components.

The rest of this paper is organized as follows. Section II introduces some essential concepts of MOPs, selection MOHHs as well as learning automata and provides background for vehicle crashworthiness. Section III presents the details of the proposed method which embeds three novel components. First, the learning automaton component designed for multiobjective optimization operates in a nontraditional way as explained in Section III-B. The second component, as described in Section III-C supports the development of a two-stage metaheuristic selection approach based on the information obtained from the learning process, enabling the use of two different metaheuristic selection methods at different stages. The third component as described in Section III-D adaptively decides when to switch to another MOEA depending on a tuned improvement threshold parameter. The parameter tuning and setting are included in Section IV, as well as the discussion and analysis of the experimental results. Section V concludes this paper and provides directions for future work.

II. BACKGROUND

A. Related Work on Multiobjective Selection Hyper-Heuristics

MOEAs and other multiobjective approaches aim to identify true Pareto fronts (PFs), i.e., equal quality optimal tradeoff solutions. If the true PFs are unknown, then MOEAs are used to generate “good” approximations [19]. The majority of the multiobjective approaches contain certain algorithmic components to achieve the following key goals [1]: 1) preserve nondominated solutions; 2) progress toward the true PFs; and 3) maintain a diverse set of solutions in the objective space.

WFG [20] and DTLZ [21] are two widely used test suites in the MOEA literature that provide benchmark functions with various characteristics. The comparison of different PFs obtained from different MOEAs is not trivial because multiple aspects should be considered, such as convergence (how close the final fronts to the true PFs are) and diversification (how dispersed the obtained fronts are) capabilities. There are a variety of performance indicators including the convergence indicators, such as *hypervolume* and $\epsilon+$ [19]. Hypervolume measures the size of objective space covered by the resultant front with respect to a reference point, while $\epsilon+$ is the minimum distance that a solution front needs to move in all dimensions to dominate the reference front. As for diversification, the most commonly used indicators include *spread* [16] and *generalized spread* [22] which extends *spread* to higher than two dimensions. Generalized spread is computed based on the mean Euclidean distance of any nearest pairs of neighbors in the nondominated solution set. The smaller the value, the better the spread of the resultant front. More analysis and review of various performance indicators for MOEAs can be found in [19].

Designing, implementing, and maintaining a (meta)heuristic for a particular problem is a time-consuming process requiring a certain level of expertise in both problem domain and heuristic optimization. Once implemented, application of a metaheuristic to a new problem domain or even a slight change in the problem description would often require the intervention of an expert. This is basically due to the fact that metaheuristics are often customized for a particular problem domain (benchmark). On the other hand, *hyper-heuristics* have emerged as automated general-purpose cross-domain optimization methods with reusable components which can be applied to multiple problem domains/benchmarks with the least modification [7]. Dealing with multiple problem domains and problem instances means dealing with various scales of objective values, making it extremely difficult to compare the cross-domain performances of algorithms. Which method to use for performance comparison of hyper-heuristics across multiple problem domains (distributions/benchmarks) and how the performance comparison should be done are still open issues in the hyper-heuristic research. Currently, there are two commonly used metrics in the area: 1) *formula 1 ranking* [9], [12] and 2) μ_{norm} [23], [24]. In this paper, we preferred the latter one (details are in Section IV-B) which is a more informative metric taking into account of the mean performance of algorithms using normalized performance indicator values over a

given number of trials on the instances from multiple problem domains/benchmarks.

The focus of this paper is on selection hyper-heuristics which choose and apply from a set of low level (meta)heuristics at each decision point of the search [8]. A key component in a selection hyper-heuristic is the (meta)heuristic selection method which should be capable of adapting itself depending on the situation to choose the appropriate low level (meta)heuristic at each decision point. Hence, learning is a crucial component of (meta)heuristic selection methods. Additionally, move acceptance technique is another key component of selection hyper-heuristics [25], [26], which determines whether or not newly generated solution(s) should be accepted as the input solution(s) to the next step/stage. The majority of the previous studies on selection hyper-heuristics focus on optimization of single objective problems. Still, there are a few studies on multiobjective selection hyper-heuristics investigating either the use of selection hyper-heuristics controlling multiple operators or mixing multiple multiobjective metaheuristics.

Hitomi and Selva [27] presented a selection hyper-heuristic (HH-AP) using an online learning heuristic selection method based on adaptive pursuit [28] managing five domain-specific perturbation operators. HH-AP is utilized for solving a multiobjective design problem for an Earth observation satellite system. Vázquez-Rodríguez and Petrovic [29] proposed a hyper-heuristic which mixes four different indicators, each from a well-established MOEA, including NSGA-II, SPEA2, and two IBEA variants to rank individuals for mating. An indicator gets selected depending on the associated probability for each individual and four subpopulations are constructed. Mating occurs within each subpopulation using binary tournament selection and eventually, four offspring pools are formed constituting to the new population. The indicator probabilities are maintained during the search via mixture experiments based on a statistical model. Kumari and Srinivas [30], [31] incorporated a roulette wheel-based heuristic selection mechanism [32] into their MOHH evolutionary algorithm to select low level mutation operators. Guizzo *et al.* [33] developed a hyper-heuristic based on two heuristic selection methods (choice function [14] and multiarmed bandit [34]) for choosing from multiple mutation and crossover operators during the search for the multiobjective integration and test order problems [35].

Some offline learning techniques have also been seen in recent MOHHs studies, e.g., genetic programming techniques in [36]–[40], grammatical evolution in [41] and [42], and top-down induction of decision trees in [43].

On the other hand, there are a few studies on multiobjective search methods that make use of multiple MOEAs. Vrugt and Robinson [44] proposed a multialgorithm genetically adaptive multiobjective (AMALGAM) method performing *cooperative search* using various MOEAs. AMALGAM executes all MOEAs simultaneously, each with a separate subpopulation at each step, and a pool of offspring gets generated by each MOEA. Those offspring pools from MOEAs are merged to form the new population. Afterwards, fast nondominated sorting is applied to the union of the new and previous

populations to choose the elite solutions surviving to the next generation. The size of the subpopulation for each MOEA gets updated adaptively based on the number of surviving solutions from each MOEA. The search continues until a set of termination criteria is satisfied. Maashi *et al.* [14] introduced a powerful online learning selection hyper-heuristic for multiobjective optimization, namely choice function-based MOHH (HH-CF) and managing NSGA-II, SPEA2, and MOGA [45]. The proposed choice function maintains an adaptively changing score for each low level MOEA during the search process based on two key components: 1) individual performance and 2) time elapsed since the last call of an MOEA. The former component uses four different indicators, including hypervolume, uniform distribution, ratio of nondominated individuals, and algorithm effort [46]. It is for exploitation, advocating the invocation of the most successful MOEA with the highest score repeatedly, while the other component is for exploration, giving a chance to the MOEAs which were used the least. The MOEA with the top score is always chosen and applied at each decision point. The results in [14] show that HH-CF outperforms not only the three underlying MOEAs which are executed individually, but also AMALGAM and a random hyper-heuristic on the majority cases of bi-objective WFG benchmark functions.

In this paper, we focus on online learning techniques as a part of selection MOHHs. It has already been observed that different MOEAs show strengths with respect to different metrics on different MOP domains [47]. The learning ability for detecting the best performing (meta)heuristic and/or identifying the *synergetic* (meta)heuristics [10], [48] over time is crucial to design an effective selection hyper-heuristic. Hence, it is reasonable to incorporate different MOEAs within an online learning selection hyper-heuristic framework for improving the cross-domain performance of the overall approach which can benefit from adaptively switching between those MOEAs over time.

HH-CF [14] is one of the best performing online learning MOHHs, to the extent of our knowledge. Similar to HH-CF, the proposed hyper-heuristics can also perform exploration and exploitation. A major difference is that the online learning method is based on *learning automata* within our selection hyper-heuristics for multiobjective optimization. Additionally, there is an adaptive mechanism to ensure that the balance between the exploration and exploitation is maintained based on the information gathered by this machine learning technique during the search. A variant of learning automata was embedded into a single-objective hyper-heuristic, i.e., AdaphH [48] which won the CHeSC competition¹ across six problem domains: 1) max-SAT; 2) bin packing; 3) personnel scheduling; 4) flow shop; 5) traveling salesman problem; and 6) vehicle routing problem. The importance of learning in selection hyper-heuristics and the success of AdaphH in solving single objective optimization problems motivated us to employ an online learning mechanism within our MOHHs for cross-domain search.

¹<http://www.asap.cs.nott.ac.uk/external/chesc2011/>

B. Learning Automata

Learning automata, introduced by Tsetlin [49] as a reinforcement learning method, has been used in a range of fields, including pattern classification [50] and signal processing [51]. A learning automaton performs an action and then classifies it as desirable or not based on a reinforcement signal (negative/penalty or positive/reward) from the environment [52]. The learning scheme then updates the reward or penalty on this action depending on the reinforcement signal. The set of actions processed by learning automata is problem dependent and varies from one application to another, for example, it could be choices of a parameter value in [50], heuristics in [51], or partitions in [53].

More formally, a learning automaton is defined as a quadruple $(A, \beta, p, \text{and } U)$, where A is the action set, β (equals to 0 or 1) represents the (penalty or reward) *feedback* or *reinforcement* signal obtained from the environment after taking the chosen action a_i at a given time t , p is the (action) selection probability vector, where each entry indicates the probability of an action being selected, and U is the update scheme. The action set A is commonly considered to be a finite set, i.e., $A = \{a_1, a_2, \dots, a_r\}$. Thus, the traditional model of a learning automaton is referred to *finite action learning automaton* [52], which is denoted as LA in this paper. At a given time (t), the action selection method chooses an action (say, a_i) based on p . After the selected action a_i is performed, p is updated by the scheme U as defined in (1) and (2) using the feedback $\beta(t)$ received from the environment. The sum of all selection probabilities in p is always equal to 1.

If a_i is the action chosen at time step t

$$p_i(t+1) = p_i(t) + \lambda^{(1)}\beta(t)(1 - p_i(t)) - \lambda^{(2)}(1 - \beta(t))p_i(t). \quad (1)$$

For other actions $a_j \neq a_i$

$$p_j(t+1) = p_j(t) - \lambda^{(1)}\beta(t)p_j(t) + \lambda^{(2)}(1 - \beta(t))\left[\frac{1}{r-1} - p_j(t)\right]. \quad (2)$$

The parameters $\lambda^{(1)}$ and $\lambda^{(2)}$ are the reward and penalty rates, respectively. When $\lambda^{(1)} = \lambda^{(2)}$, the model is referred as linear reward-penalty (L_{R-P}). In case of $\lambda^{(2)} = 0$, it is referred to as linear reward-inaction (L_{R-I}). If $\lambda^{(2)} < \lambda^{(1)}$, it is called linear reward- ϵ -penalty ($L_{R-\epsilon P}$).

C. Vehicle Crashworthiness Problem

In the automotive industry, *crashworthiness* refers to the ability of a vehicle and its components to protect its occupants during an impact or crash [54]. The crashworthiness design of vehicles is of special importance, yet, highly demanding for high-quality and low-cost industrial products. The structural optimization of the vehicle design involves multiple criteria to be considered. Liao *et al.* [55] presented a multiobjective model for the vehicle design which minimizes three objectives: 1) **weight** (mass); 2) **acceleration** characteristics (A_{in}); and 3) **toe-board intrusion** (intrusion). More specifically, the weight of the vehicle is to be minimized for enabling economic

mass production. An important goal of the vehicle design is to reduce any potential harm to occupant(s). When the front of a vehicle hits an object, it first begins to decelerate by the impact. The velocity decreases to zero when the vehicle comes to a halt. As the vehicle begins to bounce back, the velocity increases. This acceleration can cause head injuries to occupant(s) and be dangerous to other road users, because the vehicle is now moving in the opposite direction. To reduce the acceleration due to collision and possible head injuries to occupants caused by the worst scenario of the acceleration pulse [56], minimizing an integration of collision acceleration between 0.05–0.07 s in the “full frontal crash” is set as the second objective. Another mechanical injury to occupants may come from the toe-board intrusion during the crash. It could hurt the knee trajectories of occupants and influence the steering of the vehicle. Therefore, minimizing the toe-board intrusion in the 40% offset frontal crash is chosen as the third objective. **The decision variables are the thickness of five predefined reinforced points, say x_1 – x_5 , around the frontal structure of a vehicle. Each decision variable is between 1 and 3 mm.** The vehicle crashworthiness problem (VCP) model is formulated as follows:

$$\begin{aligned} &\text{minimize } F(X) = [\text{Mass}, A_{in}, \text{Intrusion}] \\ &\text{subject to } 1.0 \leq x_i \leq 3.0, \quad i = 1, 2, \dots, 5 \\ &\quad X = (x_1, x_2, \dots, x_5)^T \end{aligned} \quad (3)$$

where

$$\begin{aligned} \text{Mass} = &1640.2823 + 2.3573285x_1 + 2.3220035x_2 \\ &+ 4.5688768x_3 + 7.7213633x_4 + 4.4559504x_5 \end{aligned} \quad (4)$$

$$\begin{aligned} A_{in} = &6.5856 + 1.15x_1 - 1.0427x_2 + 0.9738x_3 \\ &+ 0.8364x_4 - 0.3695x_1x_4 + 0.0861x_1x_5 \\ &+ 0.3628x_2x_4 - 0.1106x_1^2 - 0.3437x_3^2 \\ &+ 0.1764x_4^2 \end{aligned} \quad (5)$$

$$\begin{aligned} \text{Intrusion} = &-0.0551 + 0.0181x_1 + 0.1024x_2 \\ &+ 0.0421x_3 - 0.0073x_1x_2 + 0.024x_2x_3 \\ &- 0.0118x_2x_4 - 0.0204x_3x_4 - 0.008x_3x_5 \\ &- 0.0241x_2^2 + 0.0109x_4^2. \end{aligned} \quad (6)$$

Apart from the original problem instance requiring optimization of all the three objectives, we formed additional instances by considering pairs of objectives leading to four VCPs, including VC1: minimize {Mass, A_{in} , Intrusion}, VC2: minimize {Mass, A_{in} }, VC3: minimize {Mass, Intrusion}, and VC4: minimize { A_{in} , Intrusion} for this paper.

III. METHODOLOGY

The proposed learning automata-based MOHH framework enabling control of multiple MOEAs operates as illustrated in Algorithm 1. First, given a set of MOEAs (H), the initialization process takes place to set up the relevant data structures (line 1). Our learning automaton requires the maintenance of a *transition matrix* (P) which describes the selection probabilities of metaheuristics transitioning from the previously

Algorithm 1: Learning Automaton-Based Hyper-Heuristic Framework

```

Popcurr : set (population) of input solutions,
Popnext : set of solutions surviving to the next stage,
H: set of metaheuristics (MOEAs)  $\{h_1, \dots, h_i, \dots, h_r\}$ ,
P: transition matrix, g: fixed number of generations
1 [A, P, hi, Popcurr]  $\leftarrow$  Initialise(H) ;           //  $A \subseteq H$ 
2 while (termination criteria not satisfied) do
3   Popnext  $\leftarrow$  ApplyMetaheuristic(hi, Popcurr, g);
4   Popcurr  $\leftarrow$  Replace(Popcurr, Popnext);
   // Decide whether to switch to
   // another metaheuristic
5   if (switch()) then
6     LearningAutomataUpdateScheme(P);
     // Decision Point for
     // metaheuristic selection
7     hi  $\leftarrow$  SelectMetaheuristic(P, A);
   end
end
  
```

selected metaheuristics. At the end of initialization step, the transition matrix is set up and a subset or full set of MOEAs (*A*) is determined as the input of the following learning scheme, as well as the input heuristic (*h_i*) and population (*Pop_{curr}*) (see Section III-A).

The chosen MOEA (*h_i*) is applied (line 3) to the incumbent set of solutions (*Pop_{curr}*) to the problem instance dealt with for a fixed number of generations/iterations (*g*), producing a new set of solutions (*Pop_{next}*). The new population then replaces the current population (line 4). If the conditions of switching to another metaheuristic (line 5) are satisfied, the reinforcement learning scheme updates the transition matrix (line 6) based on the feedback received during the search. Afterwards, the selection mechanism makes use of the updated transition matrix (*P*) to decide which MOEA (*h_i*) to run in the next iteration. Then all those steps are repeated until the termination criteria are satisfied.

The framework consists of four key components: 1) initialization process; 2) reinforcement learning scheme; 3) metaheuristic (action) selection method; and 4) the method deciding when to switch to another metaheuristic. Two MOHs, referred to as HH-LA and HH-RILA are designed under this framework in this paper. HH-LA and HH-RILA differ only in their initialization processes. The remaining components are the same. The following sections describe each component in detail.

A. Initialization

HH-LA utilizes all *r* MOEAs and the transition matrix *P* is initially created so that each MOEA has the same probability of being selected, i.e., $1/r$. The initial population for HH-LA is generated randomly.

HH-RILA uses a more elaborate initialization process. We propose a ranking scheme to form a reduced subset of MOEAs, eliminating the ones with relatively poor

performance. The ranking process begins with running each MOEA successively for a number of stages. The number of stages is set to the number of low level metaheuristics for giving each MOEA an equal chance to show its performance. Initial population is generated randomly. The resultant population obtained at the end of each stage is directly fed into the following stage for each MOEA. The hypervolume values for all resultant populations obtained at the end of each stage from each MOEA is computed based on the normalized objective values, i.e., $(f_i(x) - f_i^{\min}) / (f_i^{\max} - f_i^{\min})$ for the *i*th objective, where the extreme objective values for each dimension, i.e., f_i^{\max} , f_i^{\min} are updated using the maximum and minimum values found so-far by all MOEAs. This process enables performance comparison of all MOEAs with respect to hypervolume for all stages. Then we count the number of stages (frequencies), denoted as $\text{Frq}_{\text{best}}(h_i)$ (the higher, the better) that each MOEA becomes the best performing algorithm out of all stages. These counts are then used for ranking all MOEAs. If more than one MOEA has the same rank, ties are broken using the diversification indicator of generalized spread (the smaller, the better). Then MOEA(s) that rank worse than the median MOEA get excluded from the low level MOEA set. For example, if *h₃* becomes the top ranking metaheuristic in all three stages, while *h₁* and *h₂* do not in any of the three stages, then $\text{Frq}_{\text{best}}(h_1)$, $\text{Frq}_{\text{best}}(h_2)$, and $\text{Frq}_{\text{best}}(h_3)$ are 0, 0, and 3, respectively. Consequently, the rank of each MOEA with respect to normalized hypervolume is 2, 2, and 1. Suppose *h₁* has a smaller generalized spread value than *h₂*, then final ranks of *h₁*, *h₂*, and *h₃* are 2, 3, and 1, respectively. Eventually, *h₂* gets excluded from the following stage of the learning process.

Then HH-RILA operates as HH-LA with a reduced subset of low level MOEAs for the remaining search process using the final population from the best ranking MOEA as input.

B. Reinforcement Learning Scheme

The reinforcement learning scheme sits at the core of the metaheuristic selection process. The system learns a mapping (or policy) from situations to actions through a trial-and-error process with the goal of maximizing the overall reward. To explore the possible cooperation among different action pairs, the learning scheme in this paper updates the **transition probability** ($p_{(i,j)}$) **from a preceding action** (*a_i*) **to a given successor** (*a_j*), depending on **the performance after applying** *a_j*. The chosen heuristics logically form a chain of a heuristic sequence as the search progresses. Although there are previous studies [9], [48], [57]–[59] using some notion of transition probabilities to keep track of the performance of heuristics invoked successively, none of them employed the same reinforcement learning scheme as we proposed. More importantly, all the previously mentioned algorithms were tested on single objective optimization problems under a single point-based search framework managing move operators rather than metaheuristics.

In the proposed learning scheme, an action (say *h_i*) corresponds to the selection of an MOEA, and the *t*th time step is analogous to the *t*th decision point when an MOEA is

selected and applied to the tradeoff solutions in hand. The **linear reward-penalty** scheme is used to **update the transition probability from h_i to h_j at time $(t+1)$** , i.e., $p_{(i,j)}(t+1)$. The update is performed as provided in (7) and (8) [52]. The value of $\beta(t)$ is set to 1 for positive (or preferable) feedback, 0 otherwise.

If the successor metaheuristic h_j of h_i is selected

$$p_{(i,j)}(t+1) = p_{(i,j)}(t) + \lambda_{(i,j)}(t)\beta(t)(1 - p_{(i,j)}(t)) - \lambda_{(i,j)}(t)(1 - \beta(t))p_{(i,j)}(t). \quad (7)$$

For the rest of the metaheuristics that are not chosen, indexed as l , where $l \neq j$

$$p_{(i,l)}(t+1) = p_{(i,l)}(t) - \lambda_{(i,l)}(t)\beta(t)p_{(i,l)}(t) + \lambda_{(i,l)}(t)(1 - \beta(t))\left[\frac{1}{r-1} - p_{(i,l)}(t)\right]. \quad (8)$$

We use the **“change in the hypervolume value”** measured before and after selecting and applying an MOEA for reward/penalising during the learning process for two reasons. First, hypervolume is the only known unary *Pareto compliant* indicator [19], [60], i.e., if a PF P_1 dominates P_2 , the indicator value of P_1 should be better than that of P_2 . Second, theoretical studies show that maximizing the hypervolume indicator during the search is equivalent to optimizing the overall objective leading to an optimal approximation of the true PF [61], [62].

Due to the nonstationary nature of the search process, it is reasonable to **give more weight to the recent rewards than the long-past ones**. One of the common ways of doing this is to discount the past reward at a fixed ratio (α) [63]. The reward is denoted as $Q_{(i,j)}(k+1)$, meaning the estimated action value of the transition pair (h_i, h_j) occurring its $(k+1)$ th times at the t th decision point

$$Q_{(i,j)}(k+1) = Q_{(i,j)}(k) + \alpha[r_{(i,j)}(k+1) - Q_{(i,j)}(k)] \quad (9)$$

where $r_{(i,j)}(k+1)$ is the current reward obtained by pair (h_i, h_j)

$$r_{(i,j)}(k+1) = v_j(t) - v_i(t-1) \quad (10)$$

where $v_j(t)$ is the hypervolume obtained by executing the action h_j at the current t th decision point and $v_i(t-1)$ is the hypervolume obtained by action h_i at the $(t-1)$ th decision point. α is commonly fixed as 0.1 [63] as in this paper. The hypervolume here is computed in the normalized objective space as described in Section III-A.

Given the varying performance of each MOEA pair (h_i, h_j) during the search, instead of fixing the reward and penalty rates of $\lambda_{(i,j)}$, it is adaptively updated using the estimated action value of each transition pair $(Q_{(i,j)})$ at each decision point. The calculation of λ is used to update both reward and penalty rates as follows:

$$\lambda_{(i,j)}(t) = 0.1 + mQ_{(i,j)}(k+1) \quad (11)$$

where m is fixed as small positive multipliers (e.g., 2) to amplify the effect of the estimated action value $Q_{(i,j)}(t)$ on the reward/penalty parameter.

Due to the nature of the search space and amplifying multipliers, it is possible that the adaptive reward and penalty

rates ($\lambda_{(i,j)}(t)$) can get out of the $[0, 1]$ range and so the transition probabilities. In such cases, the value of $\lambda_{(i,j)}(t)$ is reset to the closest extreme value (0 or 1) ensuring that it stays within the range.

C. Metaheuristic Selection Method

In reinforcement learning, in order to take an action (i.e., choosing a metaheuristic), a selection method is required. This method is normally based on a function of the selection probabilities (utility values) to select an action at a given certain point. Several selection methods are commonly used in the scientific literature, such as roulette wheel, or greedy [63]. Those methods differ when exploring new actions and exploiting the knowledge obtained from the previous actions. The **roulette wheel** selection method chooses an action with a probability proportional to its utility value. The advantage of this method is its straightforwardness and it does not introduce any extra parameters. However, it has less chance to exploit the best-so-far actions when compared to the other selection methods, in particular when the selection probabilities of actions are similar. The **greedy selection** method only chooses the action with the highest selection probability. As a drawback, this method could overlook the other potentially good performing actions which might give higher rewards in the later stages. Further details on different selection methods can be found in [63]. Each selection method has its strengths and weaknesses. To exploit the merits of both roulette and greedy selection methods, we propose **a new selection method**, named as **ϵ -RouletteGreedy selection**. The main idea is that the selection method first focuses on exploring different transition pairs by performing a certain number of trials to get a better view of the pairwise performances of metaheuristics at the early stage. Then, the selection method becomes more and more greedy exploiting the accumulated knowledge.

The proposed selection method works as follows. The exploration phase parameter τ is fixed to a value in $[0, 1]$. During the first $\tau n_{\text{totalIter}}$ iterations, where $n_{\text{totalIter}}$ is the total number of iterations, **roulette wheel selection** is solely used to choose an action (say, h_j) out of all the possible successors of action h_i based on the transition probability $p_{(i,j)}$. Following this exploration phase, the probability ϵ of applying the **greedy selection** method is increased linearly by the formula $\tau + (1.0 - \tau)n_{\text{iter}}/n_{\text{totalIter}}$, where n_{iter} denotes the number of iterations has passed since the beginning of the algorithm. We randomly generate a value between 0 and 1. If that value is less than or equal to ϵ , the best action (with the **highest transition probability** from the previously selected action) is chosen to be performed at the next decision point. If the random value is greater than ϵ , the next action is selected by **roulette wheel selection** method.

D. Switching to Another Metaheuristic

In this paper, we propose a **threshold method** to stop the application of a selected MOEA (h_i) repeatedly, enabling the hyper-heuristic to switch to another MOEA, adaptively. A selected metaheuristic is applied as long as there is an improvement in the hypervolume as compared to the previous

iteration above an expected level. Hence, application of the selected MOEA halts if the **hypervolume improvement** $\delta(v_{\text{iter}})$ is **less than a threshold value** of Δ_v at a given iteration, or **the maximum number of iterations** (denoted as K) for applying a **low level MOEA** is exceeded. The hypervolume improvement (change) $\delta(v_{\text{iter}})$ is computed as $(v_{\text{iter}} - v_{(\text{iter}-1)})/v_{(\text{iter}-1)}$, where v_{iter} is the hypervolume of the tradeoff solutions obtained after the application of h_i at the current iteration, and $v_{(\text{iter}-1)}$ is the hypervolume obtained from h_i at the previous iteration.

IV. COMPUTATIONAL EXPERIMENTS

The proposed multiobjective selection hyper-heuristics, HH-LA and HH-RILA controlling three low level MOEAs **{NSGA-II, SPEA2, and IBEA}** are studied using a range of three-objective benchmark functions from the WFG [20] and DTLZ [21] test suites. The number of stages in the initialization for HH-RILA is set to 3. The performances of HH-LA and HH-RILA are not only compared to each individual low level MOEA, but also to random choice hyper-heuristic (HH-RC) serving as a reference approach utilizing no learning as well as the online learning hyper-heuristic of HH-CF [14] using the same set of low level MOEAs. The jMetal software platform [64] embedding implementations of the WFG and DTLZ problems and three low level MOEAs are used for the development of all the algorithms experimented within this paper.

A. Experimental Settings

Each experiment with an algorithm is repeated for 30 times on each problem instance. The WFG and DTLZ benchmark functions are all parameterized. Each WFG benchmark function has 20 distance and four position (total 24) parameters, while DTLZ1, DTLZ2–DTLZ6, and DTLZ7 have 7, 12, and 22 parameters, respectively. Those parameter values are fixed as in [20] for the WFG and [21] for the DTLZ problems.

It is commonly known that the performance of meta-heuristics can be improved through *parameter tuning*, that is, detecting the best settings (configuration) for the algorithmic parameters [65], [66]. Considering the large set of parameters and their values associated with the proposed hyper-heuristics and MOEAs used in this paper, it is not feasible to test all the combinations of settings considering the immense amount of required computational budget. Instead, parameters of HH-LA and HH-RILA are tuned based on the Taguchi experimental design [67]. Whereas, the recommended configurations and parameter settings are used for all the other algorithms, including MOEAs [16], [17], [68], [69] and HH-CF [14] from the scientific literature. Simulated binary crossover and polynomial mutation [70] are used as the MOEA operators. The distribution parameters of the crossover and mutation operators are fixed as $\{\eta_c = 20.0\}$ and $\{\eta_m = 20.0\}$, respectively. The crossover and mutation probabilities are set to $\{p_c = 0.9\}$ and $\{p_m = 1/n_p\}$, where n_p is the number of parameters. Parents are selected using the binary tournament operator [71]. The maximum number of solution evaluations for each WFG and DTLZ problem is set to 50 000 and 100 000, respectively [68].

This particular setting is always maintained for all algorithms tested in this paper for a fair performance comparison between them. The population and archive sizes are both fixed as 100 for MOEAs. The number of iterations for HH-CF and HH-RC is set to the recommended value of 25, and intensification parameter of HH-CF to 100 [14]. The number of generations for each iteration is fixed as $g = 10$ for HH-LA and HH-RILA.

For a fair comparison, the number of evaluations used for the initialization in HH-RILA are deducted from the total. As mentioned above, parameters of the proposed hyper-heuristics are tuned for an improved performance. The parameter tuning experiments and sensitivity analysis of each parameter for HH-LA and HH-RILA are provided in the following section.

B. Parameter Tuning of HH-LA and HH-RILA and Sensitivity Analysis

Our multiobjective selection hyper-heuristics contains four main parameters: 1) exploration phase τ ; 2) reward/penalty multiplier m ; 3) maximum number of iterations K for applying a low level MOEA; and 4) hypervolume improvement threshold Δ_v . Five different values for each parameter are considered: $\tau \in \{0.1, 0.3, 0.5, 0.7, 0.9\}$, $m \in \{1.0, 1.5, 2.0, 2.5, 3.0\}$, $K \in \{1, 2, 3, 4, 5\}$, and $\Delta_v \in \{0.0, 0.0025, 0.005, 0.0075, 0.01\}$. Even with this sample of five settings for each of the four parameters, 625 parameter tuning experiments would have been required for testing all combinations of the parameter settings. In this paper, the Taguchi orthogonal arrays experimental design method [67], [72] is used for parameter tuning. Sampling the configurations based on the orthogonal array, denoted as L_{25} , reduces the number of parameter tuning experiments to 25 configurations for each algorithm, which are tested on the benchmark functions.

The measurement used during the tuning experiments is μ_{norm} . The original μ_{norm} is defined for the minimization problems. Since we are maximizing hypervolume, we slightly modify the formulation of μ_{norm} as follows. Let $S_{(x,n)}$ be the set of hypervolume (30 hypervolume values in our case resulting from 30 trials) obtained by an algorithm x , where $x \in X$ on a problem n , where $n \in N$; X and N are the sets of algorithms and problems, respectively. Let $S_n^{\min} = \text{MIN}_{s \in S_{(x,n)}, \forall x \in X}$ be the minimum and $S_n^{\max} = \text{MAX}_{s \in S_{(x,n)}, \forall x \in X}$ be the maximum hypervolume obtained by all the algorithms on a problem n . The normalized hypervolume of an algorithm x on a problem n is computed as $f_{(x,n)}^{\text{norm}} = ([S_n^{\max} - \text{AVG}_{s \in S_{(x,n)}}(s)]/[S_n^{\max} - S_n^{\min}])$. The average of $f_{(x,n)}^{\text{norm}}$ defined as $\mu_{\text{norm}}(x) = \text{AVG}_{n \in N}(f_{(x,n)}^{\text{norm}})$ serves as the measurement for the tuning experiments. The lower the $\mu_{\text{norm}}(x)$ value, the better the performance of the algorithm x .

The main effects plots in Fig. 1 indicate the mean effect of each parameter setting on the performances of HH-LA and HH-RILA. The parameter setting that achieves the lowest mean μ_{norm} averaged across all trials using that setting regardless of the remaining parameter settings would be the best value for that parameter. Thus, the best configuration for HH-LA is $\{\tau = 0.5, m = 2.5, K = 3, \Delta_v = 0.0075\}$, and for HH-RILA is $\{\tau = 0.9, m = 3.0, K = 3, \Delta_v = 0.0075\}$. Both settings are used in this paper for the rest of the experiments.

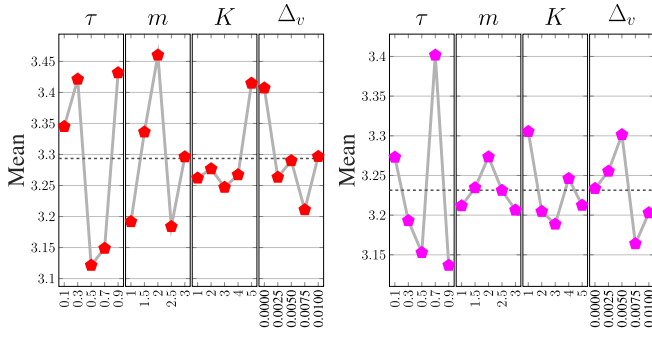


Fig. 1. Main effects plots for HH-LA (left) and HH-RILA (right) for each parameter: exploration phase (τ), multiplier (m), maximum iterations (K) for applying a low level MOEA, and hypervolume improvement threshold (Δ_v).

TABLE I

ANOVA TEST TO IDENTIFY THE CONTRIBUTION (%) OF EACH PARAMETER FOR HH-LA AND HH-RILA (DOF: DEGREES OF FREEDOM, SS: SUM OF SQUARES, MS: MEAN SQUARES, AND F: VARIANCE RATIO)

HH-LA

Parameters	DoF	SS	MS	F	p-value	contribution (%)
τ	4	0.44	0.11	8.16	0.0063	43.93
m	4	0.26	0.06	4.78	0.0289	25.75
K	4	0.09	0.02	1.73	0.2357	9.32
Δ_v	4	0.10	0.03	1.90	0.2037	10.24
Residual	8	0.11	0.01			
Total	24					100

HH-RILA

Parameters	DoF	SS	MS	F	p-value	contribution (%)
τ	4	0.24	0.06	5.05	0.0250	53.61
m	4	0.01	0.00	0.30	0.8715	3.16
K	4	0.04	0.01	0.92	0.4983	9.74
Δ_v	4	0.05	0.01	1.16	0.3972	12.27
Residual	8	0.09	0.01			
Total	24					100

Analysis of variance (ANOVA) [73] test is performed to observe how sensitive the performance of proposed hyper-heuristics to the parametric settings is by looking into the significance and contribution (in percentage) of each parameter. Table I shows that exploration phase parameter τ has the most significant influence on the performance of both HH-LA and HH-RILA at a significance level of 5% (i.e., p -value < 0.05). The parameter τ has the highest percentage contribution of 43.93% and 53.61% to the performance of HH-LA and HH-RILA, respectively. The reward/penalty multiplier m also significantly contribute to the performance of HH-LA with the second largest percentage contribution of 25.75%, while this parameter has almost no contribution (3.16%) to the performance of HH-RILA. The remaining two parameters are not significantly influential on the performance of either proposed hyper-heuristics.

C. Experimental Results on WFG and DTLZ

In this section, we use hypervolume as the main performance indicator. One-tailed Wilcoxon rank-sum test (also known as Mann–Whitney U test) is applied based on the raw hypervolume values obtained from 30 trials of each algorithm to test if there is a statistically significant performance

difference between a pair of algorithms. The significance level is set to 5%. The reference (or nadir) point (denoted as r) for the WFG and DTLZ benchmark problems are chosen as follows. For each WFG problem, the reference point is set as $r^i = 2i + 1$, where $i = 1, 2, \dots, k$ is the index of the objective and k is the total objective number. Thus, for each WFG problem, the reference point is (3, 5, 7). The reference point for DTLZ problems is set as $r^i = 0.5$ for DTLZ1, $r^i = 1.0$ for DTLZ2–DTLZ6, and $r^i = 1.0$ if $i < k$, otherwise, $r^k = 2k$ for DTLZ7.

The convergence indicator $\epsilon+$ is utilized as an additional performance comparison indicator. We notice that in some cases, the performance differences between algorithms are not distinguishable if the raw values are plotted directly. Here only for the visualization purposes, we map the raw hypervolume/ $\epsilon+$ value into the range of [0, 1] via normalization using the extreme (minimum and maximum) values collected from all algorithms over 30 trials on each instance, then the mean hypervolume and $\epsilon+$ values are plotted in Fig. 2. Higher the hypervolume or lower the $\epsilon+$ value means a better performance.

Fig. 2 shows that IBEA performs the best on WFG benchmark with respect to both hypervolume and $\epsilon+$ in the overall. HH-RILA and HH-LA follow the performance of IBEA closely. NSGA-II clearly performs the worst on WFG. The performance of IBEA gets much poorer and becomes overall the worst approach for the DTLZ benchmark functions with respect to both metrics. SPEA2 performs the best on over half of DTLZ benchmark. HH-RILA and HH-LA always achieve the second best performance on most DTLZ benchmark or even the best on DTLZ7. In addition, the hypervolume-based performance ranking of all algorithms on each benchmark problem is almost fully consistent with the $\epsilon+$ -based ranking except for WFG1. On WFG1, IBEA achieves the best rank with respect to hypervolume; however, IBEA performs slightly worse than SPEA2 on WFG1 with respect to the $\epsilon+$ indicator. This inconsistency, also discussed in [60], is possibly due to the different working principles of both indicators.

One-tailed Wilcoxon rank-sum test at 5% significance level is conducted on the performance of each pair of algorithms with respect to hypervolume. The statistical test results are summarized in Table II and we have the following observations.

In the overall, both of our MOHs deliver a better performance than any of the individual MOEAs run on its own on the WFG and DTLZ benchmarks. The statistical test results show that HH-LA and HH-RILA outperform NSGA-II on all nine WFG benchmark functions while three out of seven DTLZ problems, including DTLZ1, DTLZ2, and DTLZ5. HH-RILA additionally performs significantly better than NSGA-II on DTLZ6 and DTLZ7. HH-LA and HH-RILA perform significantly better than SPEA2 on the same eight out of nine WFG benchmark functions including WFG1 and WFG3–WFG9. HH-RILA additionally outperforms SPEA2 on DTLZ2 and DTLZ7. Although IBEA delivers a good overall performance on the WFG benchmark, both our algorithms still manage to outperform IBEA on five out of seven DTLZ problems including DTLZ1, DTLZ3, and DTLZ5–DTLZ7.

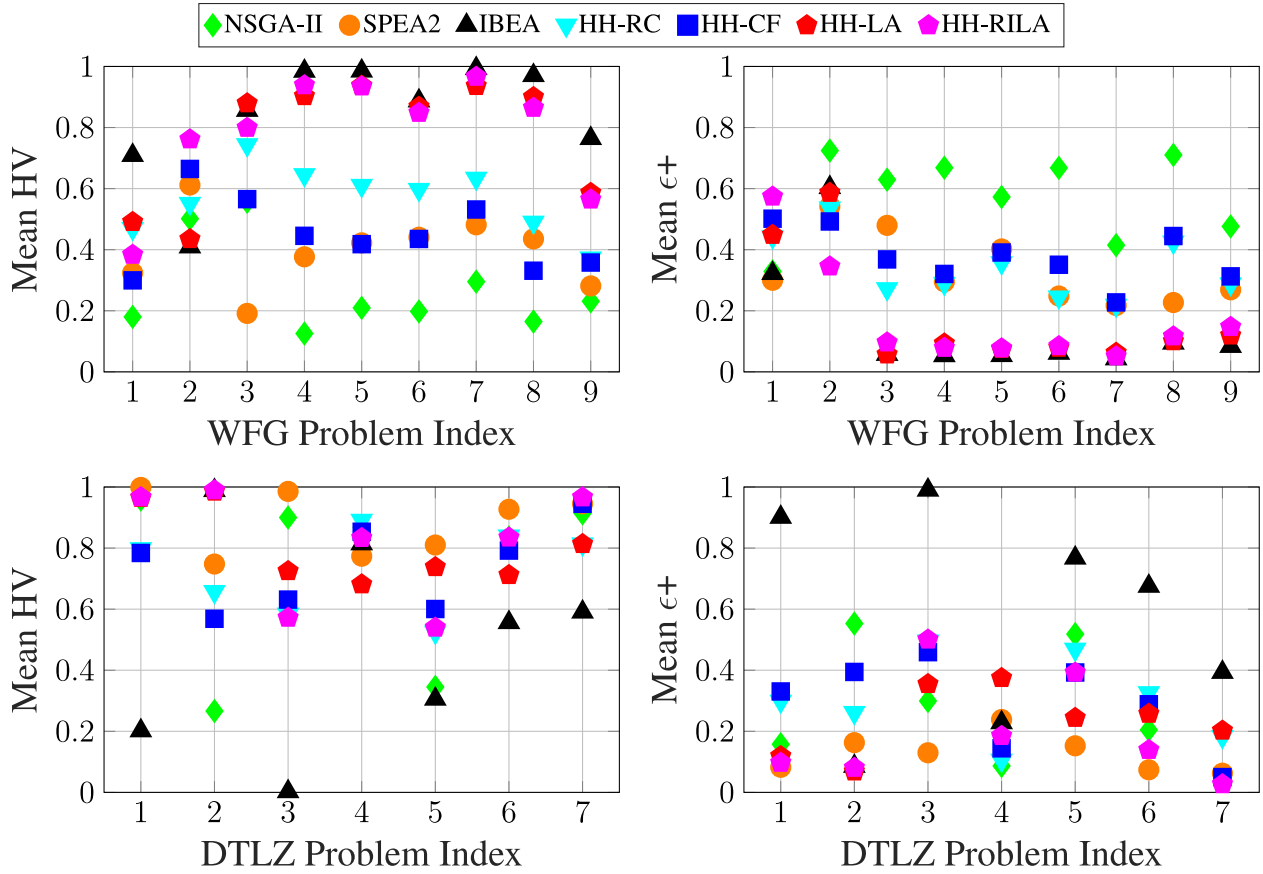


Fig. 2. Performance comparison of all the algorithms with respect to hypervolume and $\epsilon+$ on 3-D WFG and DTLZ problems.

TABLE II
ONE-TAILED WILCOXON RANK-SUM TEST AT 5% SIGNIFICANCE LEVEL ON WFG AND DTLZ BENCHMARK PROBLEMS WITH RESPECT TO HYPERVOLUME. “W” AND “D” ARE SHORT FOR “WFG” AND “DTLZ,” RESPECTIVELY. “>” MEANS THE SIGNIFICANTLY BETTER THAN, “<” SIGNIFICANTLY WORSE THAN, AND “~” NO SIGNIFICANT DIFFERENCE

	W1	W2	W3	W4	W5	W6	W7	W8	W9	D1	D2	D3	D4	D5	D6	D7
HH-LA vs HH-CF	>	~	>	>	>	>	>	>	>	>	>	>	~	~	~	~
HH-LA vs HH-RC	~	~	>	>	>	>	>	>	>	>	>	~	~	~	~	~
HH-LA vs NSGA-II	>	~	>	>	>	>	>	>	>	>	>	~	~	>	~	~
HH-LA vs SPEA2	>	~	>	>	>	>	>	>	>	>	>	>	~	>	>	>
HH-LA vs IBEA	<	~	~	<	<	~	<	<	<	<	~	~	~	~	>	>
HH-CF vs HH-RC	<	~	<	<	<	<	<	<	~	~	~	~	~	~	~	>
HH-RILA vs HH-CF	>	~	>	>	>	>	>	>	>	>	>	~	~	~	>	~
HH-RILA vs HH-RC	~	~	>	>	>	>	>	>	>	>	>	~	~	~	~	>
HH-RILA vs HH-LA	<	>	<	>	~	>	>	<	~	>	~	~	~	>	>	~
HH-RILA vs NSGA-II	>	>	>	>	>	>	>	>	>	>	>	<	~	>	>	>
HH-RILA vs SPEA2	>	~	>	>	>	>	>	>	>	>	>	<	~	>	>	>
HH-RILA vs IBEA	<	>	<	<	<	<	<	<	<	>	~	>	~	>	>	>
NSGA-II vs SPEA2	<	~	>	>	<	<	<	<	~	<	<	>	>	>	>	>
NSGA-II vs IBEA	<	~	<	<	<	<	<	<	<	<	<	>	>	~	>	>
SPEA2 vs IBEA	<	~	<	<	<	<	<	<	<	>	<	>	<	>	>	>

HH-RILA also performs significantly better than IBEA on WFG2.

Both HH-LA and HH-RILA outperform HH-CF and HH-RC. Specifically, both of our hyper-heuristics perform significantly better than HH-CF on 11 benchmark functions out of total 16, including the same eight WFG benchmark functions (WFG1 and WFG3–WFG9) and three DTLZ benchmark functions (DTLZ1–DTLZ3 for HH-LA, while DTLZ1, DTLZ2, and DTLZ6 for HH-RILA). The performance difference between each of the proposed hyper-heuristics and

HH-RC is statistically significant with respect to hypervolume on ten out of 16 problems which include the same seven WFG benchmark functions (WFG3–WFG9) and three slightly different DTLZ problems: HH-LA outperforms HH-RC on DTLZ1, DTLZ2, and DTLZ5, while DTLZ1, DTLZ2, and DTLZ7 for HH-RILA. HH-CF only outperforms HH-RC on DTLZ7, while they perform similarly on eight out of 16 problems (WFG2, WFG9, and DTLZ1–DTLZ6). HH-CF delivers a significantly worse performance than HH-RC on the rest of the seven problems (WFG1 and WFG3–WFG8).

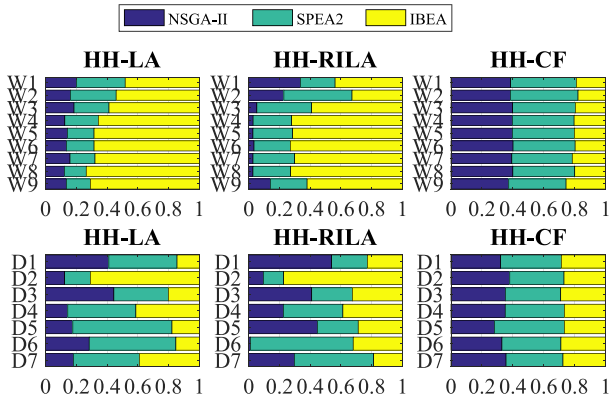


Fig. 3. Mean utilization rate of each metaheuristic by HH-LA (left), HH-RILA (middle), and HH-CF (right) over 30 trials on WFG (W) and DTLZ (D).

As for the performance comparison between HH-LA and HH-RILA, HH-LA is slightly better than HH-RILA in the overall on the WFG problems. This performance difference is statistically significant on four WFG problems including WFG1, WFG3, WFG6, and WFG8, while HH-RILA performs significantly better than HH-LA on three WFG problems: WFG2, WFG4, and WFG7. However, considering DTLZ benchmark, HH-RILA performs slightly better than HH-LA in the overall. This performance difference is statistically significant on DTLZ1 and DTLZ6, while HH-LA only outperforms HH-RILA on DTLZ5.

D. Analysis of Hyper-Heuristics on WFG and DTLZ

1) *Utilization of Low Level Metaheuristics:* The “utilization rate” of a low level metaheuristic is the number of invocations of this metaheuristic divided by the total number of metaheuristic selection decision points in a given trial. The mean utilization rates of the three MOEAs, i.e., NSGA-II, SPEA2, and IBEA averaged over 30 trials on the WFG and DTLZ benchmark functions produced by HH-LA, HH-RILA, and HH-CF [14] are illustrated in Fig. 3.

Fig. 3 shows the differences in learning characteristics of these three online learning MOHs. First, HH-LA and HH-RILA provide a bias toward using the best performing MOEA with respect to hypervolume. Specifically, both HH-LA and HH-RILA choose IBEA and SPEA2 more frequently while solving the WFG and DTLZ problems, respectively. This is not surprising, considering that hypervolume serves as the main guidance in the learning mechanisms of our hyper-heuristics. Second, in certain cases, such as WFG4–WFG8 and DTLZ6, HH-RILA almost exclude NSGA-II which is the worst performed MOEA on those problems. Interestingly, HH-CF generates a similar utilization rate for low level MOEAs across different problem sets. On average, HH-CF uses NSGA-II, SPEA2, and IBEA for 40%, 40%, and 20% of all the decision points, respectively, on the WFG benchmark. Similarly, HH-CF uses NSGA-II, SPEA2, and IBEA for 34%, 38%, and 28%, respectively, on the DTLZ benchmark. This might indicate that the adaptation mechanism in HH-CF has some issues controlling these

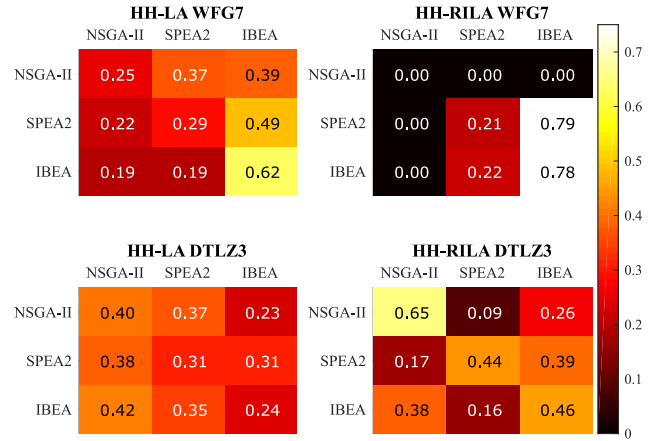


Fig. 4. Averaged transition probability matrices (over 30 trials) produced by HH-LA (left column) and HH-RILA (right column) while solving WFG7 and DTLZ3. The lighter the color, the higher the transition probability.

three low level metaheuristics properly on different problem instances.

2) *Analysis of the Transition Probabilities:* The proposed hyper-heuristics embed a learning mechanism which maintains the transition probabilities between any pair of MOEAs. Fig. 4 provides the final transition probability matrices obtained by HH-LA and HH-RILA averaged over 30 trials for the sample cases of WFG7 and DTLZ3.

Fig. 4 illustrates that both HH-LA and HH-RILA yield higher probability entries preferring transitions to IBEA than to other MOEAs for WFG7. This is consistent with the performance assessment of each individual MOEA (Fig. 2), which shows that IBEA performs the best on WFG7. Moreover, HH-RILA excludes the worst performing MOEA, i.e., NSGA-II after the initialization stage for solving WFG7. This is likely the reason why HH-RILA performs significantly better than HH-LA on WFG7.

DTLZ3 is an interesting case. IBEA delivers a better performance in the early stages, but stagnates and even deteriorates later during the search process. Due to the misleading performance of IBEA in the early stage, HH-RILA rewards IBEA more than the other MOEAs, while excluding the ones with potentially good performance, such as SPEA2. Consequently, HH-RILA ends up performing significantly worse than HH-LA on DTLZ3.

In summary, the proposed learning mechanism is capable of adaptively updating the transition probabilities between pairs of MOEAs giving bias toward the right algorithms (with good performance) during the search process in an online manner. Moreover, the ranking initialization scheme is, in some cases, capable of improving the overall performance significantly by detecting and excluding potentially poor performing MOEA(s) in the early stages of the search.

3) *Analysis of Approximate Pareto Fronts:* So far, IBEA is a strong competitor of HH-LA and HH-RILA with respect to hypervolume. To get more insights on the distribution of solutions from IBEA and the proposed hyper-heuristics, PFs obtained from HH-RILA and IBEA for WFG7 and DTLZ3 are illustrated in Fig. 5. HH-LA produces PFs very similar

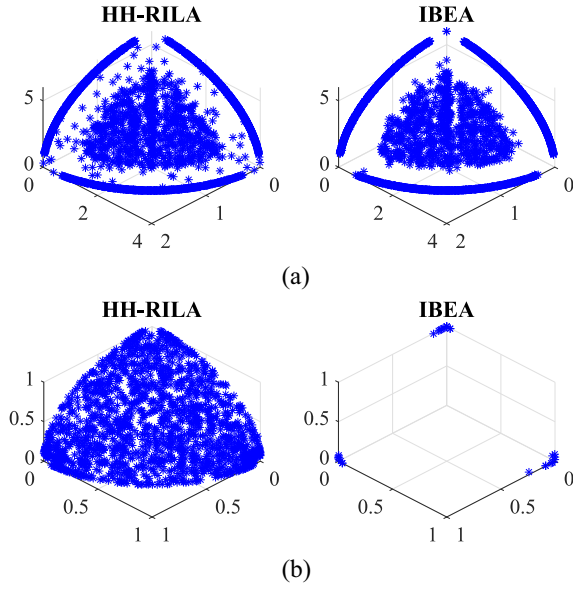


Fig. 5. Approximate PFs produced by HH-RILA (left column) and IBEA (right column) on (a) WFG7 and (b) DTLZ3.

to HH-RILA on almost all problems, and so we focus on HH-RILA here.

Fig. 5 demonstrates that IBEA is prone to be trapped at a local optimum. IBEA produces uneven solution distribution for WFG7, leaving clear gaps between the boundary and inner regions, whereas HH-RILA reaches a better solution distribution for this problem.

IBEA performs poorly on DTLZ3. All the solutions are clustered around the “corner” points which suggests that the performance of IBEA degrades during the search process. This interesting behavior of IBEA has also been observed previously by Tušar and Filipič [69, Fig. 7] and Li *et al.* [74]. More importantly, solutions from HH-RILA clearly spread much more evenly on the front than IBEA, possibly due to the utilization of multiple MOEAs.

4) *Analysis of Search Dynamics*: To have a further understanding of the search progress induced by each low level MOEA at each trial under the proposed hyper-heuristics, we recorded which MOEA was chosen and applied at each iteration of 30 trials. Fig. 6 provides plots based on that data as a representative of the search dynamics for HH-LA and HH-RILA on WFG7 and DTLZ3. Each tick on the plot indicates the total number of trials that the relevant MOEA is selected and invoked by the indicated hyper-heuristic at a given iteration. For example, HH-LA calls NSGA-II, 11 times, SPEA2, 7 times, and IBEA, 12 times at the first iteration out of 30 trials when solving WFG7.

As observed from Fig. 6, HH-LA slowly reduces the usage of poor performing MOEAs which are NSGA-II and SPEA2 for WFG7 during the exploration phase which corresponds to the first half of the search process ($\tau = 0.5$). Then the best performing MOEA; i.e., IBEA is predominately used for the remaining iterations. Meanwhile, HH-RILA excludes NSGA-II after the initial ranking stage (first nine iterations). Following the exploration phase, HH-RILA prefers invoking

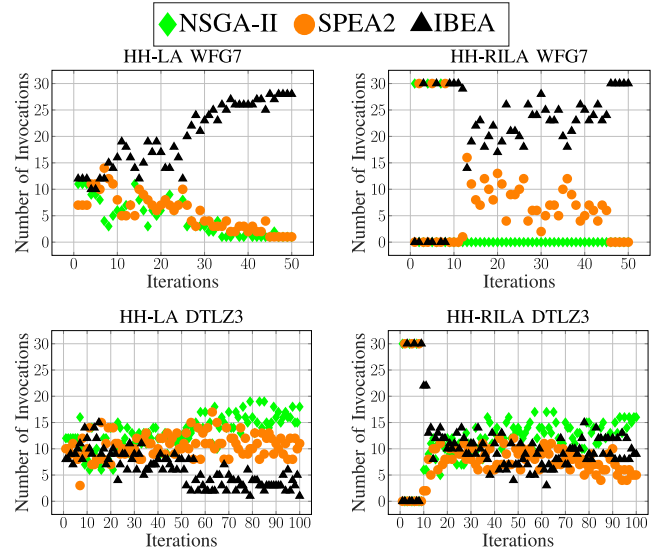


Fig. 6. Number of invocations of each MOEA at each iteration over 30 trials on WFG7 and DTLZ3 obtained from HH-LA (left column) and HH-RILA (right column).

IBEA more and more while SPEA2 less and less as the search progresses. This behavior is also reflected on the transition probability matrix (Fig. 4). The transition probabilities from any MOEA to IBEA are higher than to the other MOEAs for both HH-LA and HH-RILA. Moreover, the transition probabilities to and from NSGA-II are all 0s for HH-RILA on WFG7.

On DTLZ3, HH-LA tends to prefer employing, first, NSGA-II and then SPEA2 more than IBEA after the exploration phase. Nevertheless, HH-RILA cannot detect the degrading performance of IBEA rapidly enough during the search, and this results in excluding SPEA2 or NSGA-II in some trials. Therefore, the performance of HH-RILA becomes worse than HH-LA with respect to hypervolume. This behavior is also consistent with the mean utilization rates as provided in Fig. 3, which shows that HH-LA uses NSGA-II and SPEA2 more than HH-RILA does on DTLZ3.

E. Experimental Results for Vehicle Crashworthiness

HH-LA and HH-RILA are also tested on the real-world problem of VCP. The same experimental settings as in [14] are used for a fair performance comparison. The population size is fixed as 30 in these experiments. Each iteration consists of 50 generations. Each algorithm terminates whenever 75 000 solution evaluations are exceeded. The reference point for each VC instance is set to $r^i = z^{\text{nadir}_i} + 0.5(z^{\text{nadir}_i} - z^{\text{ideal}_i})$, where i is the index of an objective, z^{nadir_i} and z^{ideal_i} are the nadir (highest) and ideal (lowest) value of objective i , respectively. The performances of each algorithm with respect to hypervolume and $\epsilon+$ are provided in Table III. The results show that NSGA-II performs the best (with the highest hypervolume and lowest $\epsilon+$) on VC1 and VC2, while HH-RILA performs the best on VC3 and VC4 based on both performance indicators when compared to the other algorithms including each MOEA

TABLE III
MEAN_{STD} HYPERVOLUME AND $\epsilon+$ VALUES FOR THE VCPs, AVERAGED OVER 30 TRIALS AND THEIR STANDARD DEVIATIONS PROVIDED AS SUBSCRIPT ENTRIES. THE BEST MEAN HYPERVOLUME AND $\epsilon+$ ON EACH PROBLEM INSTANCE IS HIGHLIGHTED IN BOLD

	Hypervolume				$\epsilon+$			
	VC1	VC2	VC3	VC4	VC1	VC2	VC3	VC4
HH-LA	98.0693 _{1.55}	53.0715 _{2.35}	1.6553 _{0.10}	2.0558 _{0.07}	0.1354 _{0.04}	0.0717 _{0.12}	0.3565 _{0.13}	0.0844 _{0.07}
HH-RILA	98.8209 _{1.94}	53.7729 _{1.26}	1.7778_{0.11}	2.0924_{0.03}	0.1225 _{0.05}	0.0374 _{0.06}	0.1995_{0.15}	0.0495_{0.03}
HH-CF	98.9038 _{1.82}	53.5628 _{1.74}	1.6696 _{0.11}	2.0582 _{0.05}	0.1194 _{0.05}	0.0501 _{0.09}	0.3372 _{0.14}	0.0794 _{0.05}
HH-RC	98.2566 _{1.75}	52.6460 _{2.78}	1.7182 _{0.13}	2.0457 _{0.06}	0.1380 _{0.04}	0.0982 _{0.14}	0.2738 _{0.16}	0.0897 _{0.06}
NSGA-II	100.2163_{1.48}	53.9514_{0.04}	1.7093 _{0.12}	2.0880 _{0.04}	0.1007_{0.04}	0.0362_{0.00}	0.2864 _{0.16}	0.0552 _{0.04}
SPEA2	99.5532 _{0.02}	53.6523 _{1.76}	1.7278 _{0.12}	2.0669 _{0.05}	0.1188 _{0.05}	0.0486 _{0.09}	0.2651 _{0.16}	0.0818 _{0.06}
IBEA	96.9435 _{0.57}	52.1465 _{3.06}	1.6846 _{0.13}	2.0213 _{0.06}	0.1570 _{0.02}	0.1184 _{0.16}	0.3119 _{0.16}	0.1148 _{0.06}

TABLE IV

ONE-TAILED WILCOXON RANK-SUM TEST AT 5% SIGNIFICANCE LEVEL ON VC PROBLEMS WITH RESPECT TO HYPERVOLUME. > MEANS THE SIGNIFICANTLY BETTER THAN, < SIGNIFICANTLY WORSE THAN, AND ~ NO SIGNIFICANT DIFFERENCE

	VC1	VC2	VC3	VC4
HH-LA vs HH-CF	~	~	~	~
HH-LA vs HH-RC	~	>	~	~
HH-LA vs NSGA-II	<	<	~	<
HH-LA vs SPEA2	~	~	~	~
HH-LA vs IBEA	>	>	~	>
HH-RILA vs HH-CF	~	~	>	>
HH-RILA vs HH-RC	~	>	>	>
HH-RILA vs HH-LA	~	>	>	>
HH-RILA vs NSGA-II	<	<	>	~
HH-RILA vs SPEA2	~	>	~	>
HH-RILA vs IBEA	>	>	>	>
HH-CF vs HH-RC	~	~	<	~
NSGA-II vs SPEA2	~	>	<	>
NSGA-II vs IBEA	>	>	>	>
SPEA2 vs IBEA	>	>	>	>

used on its own. The proposed hyper-heuristic exploits the synergy between MOEAs with different performances leading to an improved overall performance.

The one-tailed Wilcoxon rank-sum test comparing the performance difference of all pairs of multiobjective approaches are summarized in Table IV. The results suggest that both NSGA-II and SPEA2 perform significantly better than IBEA. HH-RILA outperforms HH-CF on two out of four VC problem instances (VC3 and VC4), and HH-RC on three out of four VC problem instances (VC2–VC4), and beats IBEA on all four VC problems. HH-LA performs similar as HH-CF, significantly better HH-RC on VC2, and also outperforms IBEA on all four VC problems. HH-RILA outperforms HH-LA on VC2–VC4.

F. Generality Analysis

Different approaches perform the best on different problems. The cross-domain search performance of algorithms indicating which algorithm is more general and the best across a range of problems is often of interest [7]. As mentioned before, μ_{norm} is used in this paper to assess the generality level of an algorithm across different problems including the benchmark functions and VCPs.

The empirical results based on μ_{norm} are presented in Table V. We also rank all the algorithms with respect to μ_{norm} values. The lower μ_{norm} value is, the better the rank of an algorithm. From Table V, we have the following

TABLE V

HYPERVOLUME μ_{norm} FOR EACH ALGORITHM ON EACH TEST PROBLEM INSTANCE. \bar{x}_W , \bar{x}_D , \bar{x}_{VC} , \bar{x}_{WD} , AND \bar{x}_{All} DENOTE THE MEAN μ_{norm} ON WFG, DTLZ, VC, AND BOTH WFG AND DTLZ, AS WELL AS ALL THREE PROBLEM DOMAINS, RESPECTIVELY. THE BEST MEAN μ_{norm} ON EACH BENCHMARK IS HIGHLIGHTED IN BOLD, THE SECOND BEST IS IN GRAY BOX

	NSGA-II	SPEA2	IBEA	HH-RC	HH-CF	HH-LA	HH-RILA
W1	0.8200	0.6761	0.2911	0.5304	0.7004	0.5091	0.6165
W2	0.4985	0.3881	0.5906	0.4486	0.3356	0.5653	0.2386
W3	0.4431	0.8086	0.1435	0.2564	0.4344	0.1206	0.2008
W4	0.8744	0.6233	0.0154	0.3545	0.5546	0.0971	0.0614
W5	0.7905	0.5776	0.0148	0.3897	0.5819	0.0630	0.0653
W6	0.8021	0.5591	0.1126	0.4029	0.5648	0.1318	0.1519
W7	0.7047	0.5185	0.0066	0.3662	0.4684	0.0639	0.0337
W8	0.8353	0.5647	0.0295	0.5101	0.6688	0.1002	0.1354
W9	0.7689	0.7186	0.2357	0.6273	0.6421	0.4141	0.4357
\bar{x}_W	0.7264	0.6038	0.1600	0.4318	0.5501	0.2294	0.2155
D1	0.0411	0.0014	0.7981	0.2021	0.2160	0.0354	0.0335
D2	0.7336	0.2522	0.0114	0.3425	0.4315	0.0150	0.0107
D3	0.1000	0.0144	0.9974	0.4164	0.3687	0.2753	0.4280
D4	0.1327	0.2266	0.1859	0.1092	0.1465	0.3186	0.1669
D5	0.6549	0.1900	0.6945	0.4757	0.3996	0.2622	0.4605
D6	0.1654	0.0729	0.4437	0.1606	0.2082	0.2882	0.1654
D7	0.0862	0.0548	0.4093	0.1866	0.0550	0.1867	0.0339
\bar{x}_D	0.2734	0.1160	0.5057	0.2705	0.2608	0.1974	0.1855
VC1	0.2715	0.3746	0.7804	0.5762	0.4756	0.6054	0.4885
VC2	0.0288	0.0711	0.2839	0.2133	0.0837	0.1531	0.0540
VC3	0.4961	0.4402	0.5708	0.4693	0.6164	0.6596	0.2887
VC4	0.1291	0.2265	0.4378	0.3247	0.2672	0.2782	0.1088
\bar{x}_{VC}	0.2314	0.2781	0.5182	0.3959	0.3607	0.4241	0.2350
\bar{x}_{WD}	0.5282	0.3904	0.3113	0.3612	0.4235	0.2154	0.2024
\bar{x}_{All}	0.4688	0.3680	0.3526	0.3682	0.4110	0.2571	0.2089
Rank	7	4	3	5	6	2	1

observations. There is no single MOEA performing well across all three problem domains even though each uses the same perturbative operators. The performance of each individual MOEA varies vastly on different domains. IBEA, SPEA2, and NSGA-II deliver the best performance on WFG, DTLZ, and VC, respectively. Although HH-RILA is always the second best performing approach on each domain, it has the best cross-domain performance and so the most general approach among all algorithms tested in this paper with a mean μ_{norm} of 0.2089 based on the results from all 20 problem instances. Moreover, HH-LA turns out to be the second best approach with a mean μ_{norm} of 0.2571. The cross-domain performance of our selection hyper-heuristics are followed by IBEA, SPEA2, HH-RC, HH-CF, and NSGA-II in that order. If we consider only WFG and DTLZ benchmarks ignoring the VCP and compare the cross-domain performance of all algorithms, HH-RILA and HH-LA still rank the best and

the second best based on the mean μ_{norm} values averaged over all benchmark functions, yielding 0.2024 and 0.2154, respectively.

V. CONCLUSION

In this paper, we proposed two variants of a learning automata-based multiobjective selection hyper-heuristic: HH-LA and HH-RILA. The performance and generality of HH-LA and HH-RILA controlling three MOEAs are investigated on three multiobjective continuous optimization problem domains, including two sets of benchmark functions (WFG and DTLZ) and the real-world problem of vehicle crash-worthiness. The experimental results show that the different MOEAs perform the best on different problem domains. However, HH-RILA and HH-LA perform the best and second best, respectively, across all test problem instances based on the μ_{norm} indicator. In addition, both HH-LA and HH-RILA outperform a state-of-the-art online learning hyper-heuristic and a random choice hyper-heuristic. The results also suggest that the proposed hyper-heuristics are indeed capable of exploiting the strengths of the low level MOEAs delivering an improved performance via a learning automata in this paper. The proposed component for metaheuristic (action) selection, that is the ϵ -RouletteGreedy method adaptively balances exploration and exploitation of the use of MOEAs by combining the merits of two regular selection methods of roulette wheel and greedy selection.

In future work, both HH-LA and HH-RILA will be applied to other discrete and continuous real-world problems and their performance as well as the level of generality will be assessed. In single objective optimization, it has been observed that the performance of selection hyper-heuristics might vary depending on the set of low level heuristics [25], hence another interesting research direction would be investigating the performance of the proposed hyper-heuristics using different sets of low level multiobjective metaheuristics. Moreover, there is a growing interest into many-objective (problems with more than three objectives) optimization and recent studies show that many-objective optimization requires different approaches from multiobjective metaheuristics [75]. The tests that we performed in this paper is somewhat at the boundary of multi and many-objective optimization considering that three objective problems are used in the experiments. Therefore, the low level metaheuristics used in the proposed multiobjective selection hyper-heuristic framework can be replaced by other state-of-the-art many-objective metaheuristics (such as NSGA-III [76] and AGE [68]), and tested on a range of many-objective optimization problems. Although hypervolume is utilized in this paper as the main indicator for guiding the search and combining strengths of different MOEAs for multiobjective optimization, the proposed framework allows the use of different indicators. It would be worth investigating the behavior of the proposed selection hyper-heuristics when other convergence/performance indicators are used for guidance, such as the $\epsilon+$ indicator.

ACKNOWLEDGMENT

The authors sincerely would like to thank Dr. Imo Eyoh, Dr. Mashaël Maashi, Dr. Markus Wagner and Dr. Wil Ward for their valuable advice, feedback and support.

REFERENCES

- [1] C. A. C. Coello, D. A. Van Veldhuizen, and G. B. Lamont, *Evolutionary Algorithms for Solving Multi-Objective Problems*, vol. 242. New York, NY, USA: Springer, 2002.
- [2] M. Saadateseresh, A. Mansourian, and M. Taleai, "Evacuation planning using multiobjective evolutionary optimization approach," *Eur. J. Oper. Res.*, vol. 198, no. 1, pp. 305–314, 2009.
- [3] T. Hanne and S. Nickel, "A multiobjective evolutionary algorithm for scheduling and inspection planning in software development projects," *Eur. J. Oper. Res.*, vol. 167, no. 3, pp. 663–678, 2005.
- [4] B. Alatas, E. Akin, and A. Karci, "Modenar: Multi-objective differential evolution algorithm for mining numeric association rules," *Appl. Soft Comput.*, vol. 8, no. 1, pp. 646–656, 2008.
- [5] E. Masazade *et al.*, "A multiobjective optimization approach to obtain decision thresholds for distributed detection in wireless sensor networks," *IEEE Trans. Syst., Man, Cybern. B, Cybern.*, vol. 40, no. 2, pp. 444–457, Apr. 2010.
- [6] A. Zhou *et al.*, "Multiobjective evolutionary algorithms: A survey of the state of the art," *Swarm Evol. Comput.*, vol. 1, no. 1, pp. 32–49, 2011.
- [7] E. K. Burke *et al.*, "Hyper-heuristics: A survey of the state of the art," *J. Oper. Res. Soc.*, vol. 64, no. 12, pp. 1695–1724, Dec. 2013, doi: [10.1057/jors.2013.71](https://doi.org/10.1057/jors.2013.71).
- [8] E. K. Burke *et al.*, "A classification of hyper-heuristic approaches," in *Handbook of Metaheuristics*. Boston, MA, USA: Springer, 2010, pp. 449–468.
- [9] S. Asta, E. Özcan, and T. Curtois, "A tensor based hyper-heuristic for nurse rostering," *Knowl. Based Syst.*, vol. 98, pp. 185–199, Apr. 2016.
- [10] S. Asta and E. Özcan, "A tensor-based selection hyper-heuristic for cross-domain heuristic search," *Inf. Sci.*, vol. 299, pp. 412–432, Apr. 2015.
- [11] T. Wauters, K. Verbeeck, P. De Causmaecker, and G. V. Berghe, "Boosting metaheuristic search using reinforcement learning," in *Hybrid Metaheuristics*, vol. 434. Heidelberg, Germany: Springer, 2013, pp. 433–452.
- [12] A. Kheiri and E. Özcan, "An iterated multi-stage selection hyper-heuristic," *Eur. J. Oper. Res.*, vol. 250, no. 1, pp. 77–90, 2016.
- [13] K. Z. Zamli, F. Din, G. Kendall, and B. S. Ahmed, "An experimental study of hyper-heuristic selection and acceptance mechanism for combinatorial T-way test suite generation," *Inf. Sci.*, vol. 399, pp. 121–153, Aug. 2017.
- [14] M. Maashi, E. Özcan, and G. Kendall, "A multi-objective hyper-heuristic based on choice function," *Expert Syst. Appl.*, vol. 41, no. 9, pp. 4475–4493, 2014.
- [15] R. A. Gonçalves, J. N. Kuk, C. P. Almeida, and S. M. Venske, "MOEA/D-HH: A hyper-heuristic for multi-objective problems," in *Proc. Int. Conf. Evol. Multi Criterion Optim.*, Guimarães, Portugal, 2015, pp. 94–108.
- [16] K. Deb, A. Pratap, S. Agarwal, and T. Meyarivan, "A fast and elitist multiobjective genetic algorithm: NSGA-II," *IEEE Trans. Evol. Comput.*, vol. 6, no. 2, pp. 182–197, Apr. 2002.
- [17] E. Zitzler, M. Laumanns, and L. Thiele, "SPEA2: Improving the strength Pareto evolutionary algorithm," *ETH Zurich, Zürich, Switzerland, Rep.* 103, 2001.
- [18] E. Zitzler and S. Künzli, "Indicator-based selection in multiobjective search," in *Parallel Problem Solving From Nature—PPSN VIII*. Heidelberg, Germany: Springer, 2004, pp. 832–842.
- [19] E. Zitzler, L. Thiele, M. Laumanns, C. M. Fonseca, and V. G. Da Fonseca, "Performance assessment of multiobjective optimizers: An analysis and review," *IEEE Trans. Evol. Comput.*, vol. 7, no. 2, pp. 117–132, Apr. 2003.
- [20] S. Huband, P. Hingston, L. Barone, and L. While, "A review of multiobjective test problems and a scalable test problem toolkit," *IEEE Trans. Evol. Comput.*, vol. 10, no. 5, pp. 477–506, Oct. 2006.
- [21] K. Deb, L. Thiele, M. Laumanns, and E. Zitzler, *Scalable Test Problems for Evolutionary Multiobjective Optimization*. London, U.K.: Springer, 2005.

- [22] A. Zhou, Y. Jin, Q. Zhang, B. Sendhoff, and E. Tsang, "Combining model-based and genetics-based offspring generation for multi-objective optimization using a convergence criterion," in *Proc. IEEE Int. Conf. Evol. Comput.*, Vancouver, BC, Canada, 2006, pp. 892–899.
- [23] L. Di Gaspero and T. Urli, "Evaluation of a family of reinforcement learning cross-domain optimization heuristics," in *Learning and Intelligent Optimization*. Heidelberg, Germany: Springer, 2012, pp. 384–389.
- [24] S. Adriaensen, G. Ochoa, and A. Nowé, "A benchmark set extension and comparative study for the HyFlex framework," in *Proc. IEEE Congr. Evol. Comput. (CEC)*, Sendai, Japan, 2015, pp. 784–791.
- [25] B. Bilgin, E. Özcan, and E. E. Korkmaz, "An experimental study on hyper-heuristics and exam timetabling," in *Proc. Int. Conf. Pract. Theory Autom. Timetabling*, Brno, Czech Republic, 2006, pp. 394–412.
- [26] E. Özcan, B. Bilgin, and E. E. Korkmaz, "A comprehensive analysis of hyper-heuristics," *Intell. Data Anal.*, vol. 12, no. 1, pp. 3–23, 2008.
- [27] N. Hitomi and D. Selva, "A hyperheuristic approach to leveraging domain knowledge in multi-objective evolutionary algorithms," in *Proc. ASME Int. Design Eng. Tech. Conf. Comput. Inf. Eng. Conf.*, 2016, Art. no. V02BT03A030.
- [28] D. Thierens, "Adaptive strategies for operator allocation," in *Parameter Setting in Evolutionary Algorithms*. Heidelberg, Germany: Springer, 2007, pp. 77–90.
- [29] J. A. Vázquez-Rodríguez and S. Petrovic, "A mixture experiments multi-objective hyper-heuristic," *J. Oper. Res. Soc.*, vol. 64, no. 11, pp. 1664–1675, 2013.
- [30] A. C. Kumari and K. Srinivas, "Scheduling and inspection planning in software development projects using multi-objective hyper-heuristic evolutionary algorithm," *Int. J. Softw. Eng. Appl.*, vol. 4, no. 3, p. 45, 2013.
- [31] A. C. Kumari and K. Srinivas, "Hyper-heuristic approach for multi-objective software module clustering," *J. Syst. Softw.*, vol. 117, pp. 384–401, Jul. 2016.
- [32] A. Nareyek, "Choosing search heuristics by non-stationary reinforcement learning," in *Metaheuristics: Computer Decision-Making*. Boston, MA, USA: Springer, 2004, pp. 523–544.
- [33] G. Guizzo, G. M. Fritsche, S. R. Vergilio, and A. T. R. Pozo, "A hyper-heuristic for the multi-objective integration and test order problem," in *Proc. Annu. Conf. Genet. Evol. Comput.*, Madrid, Spain, 2015, pp. 1343–1350.
- [34] A. Fialho, L. Da Costa, M. Schoenauer, and M. Sebag, "Analyzing bandit-based adaptive operator selection mechanisms," *Ann. Math. Artif. Intell.*, vol. 60, nos. 1–2, pp. 25–64, 2010.
- [35] W. K. G. Assunção, T. E. Colanzi, S. R. Vergilio, and A. Pozo, "A multi-objective optimization approach for the integration and test order problem," *Inf. Sci.*, vol. 267, pp. 119–139, May 2014.
- [36] A. Masood, Y. Mei, G. Chen, and M. Zhang, "Many-objective genetic programming for job-shop scheduling," in *Proc. IEEE Congr. Evol. Comput. (CEC)*, Vancouver, BC, Canada, 2016, pp. 209–216.
- [37] A. Masood, Y. Mei, G. Chen, and M. Zhang, "A PSO-based reference point adaption method for genetic programming hyper-heuristic in many-objective job shop scheduling," in *Proc. Aust. Conf. Artif. Life Comput. Intell. (ACALCI)*, Melbourne, VIC, Australia, 2017, pp. 326–338.
- [38] D. Karunakaran, G. Chen, and M. Zhang, "Parallel multi-objective job shop scheduling using genetic programming," in *Proc. Aust. Conf. Artif. Life Comput. Intell.*, Canberra, ACT, Australia, 2016, pp. 234–245.
- [39] S. Nguyen, M. Zhang, M. Johnston, and K. C. Tan, "Automatic design of scheduling policies for dynamic multi-objective job shop scheduling via cooperative coevolution genetic programming," *IEEE Trans. Evol. Comput.*, vol. 18, no. 2, pp. 193–208, Apr. 2014.
- [40] M. Freitag and T. Hildebrandt, "Automatic design of scheduling rules for complex manufacturing systems by multi-objective simulation-based optimization," *CIRP Ann. Manuf. Technol.*, vol. 65, no. 1, pp. 433–436, 2016.
- [41] C. Ryan, J. J. Collins, and M. O. Neill, "Grammatical evolution: Evolving programs for an arbitrary language," in *Proc. Eur. Conf. Genet. Program.*, Paris, France, 1998, pp. 83–96.
- [42] T. Mariani, G. Guizzo, S. R. Vergilio, and A. T. R. Pozo, "Grammatical evolution for the multi-objective integration and test order problem," in *Proc. Genet. Evol. Comput. Conf.*, Denver, CO, USA, 2016, pp. 1069–1076.
- [43] M. P. Basgalupp, R. C. Barros, and V. Podgorelec, "Evolving decision-tree induction algorithms with a multi-objective hyper-heuristic," in *Proc. 30th Annu. ACM Symp. Appl. Comput.*, Salamanca, Spain, 2015, pp. 110–117.
- [44] J. A. Vrugt and B. A. Robinson, "Improved evolutionary optimization from genetically adaptive multimethod search," *Proc. Nat. Acad. Sci. USA*, vol. 104, no. 3, pp. 708–711, 2007.
- [45] C. M. Fonseca and P. J. Fleming, "Multiobjective optimization and multiple constraint handling with evolutionary algorithms. I. A unified formulation," *IEEE Trans. Syst., Man, Cybern. A, Syst., Humans*, vol. 28, no. 1, pp. 26–37, Jan. 1998.
- [46] K. C. Tan, T. H. Lee, and E. F. Khor, "Evolutionary algorithms for multi-objective optimization: Performance assessments and comparisons," *Artif. Intell. Rev.*, vol. 17, no. 4, pp. 251–290, 2002.
- [47] M. Maashi, G. Kendall, and E. Özcan, "Choice function based hyper-heuristics for multi-objective optimization," *Appl. Soft Comput.*, vol. 28, pp. 312–326, Mar. 2015.
- [48] M. Misir, K. Verbeeck, P. De Causmaecker, and G. V. Berghe, "A new hyper-heuristic implementation in HyFlex: A study on generality," in *Proc. 5th Multidiscipl. Int. Scheduling Conf. Theory Appl.*, 2011, pp. 374–393.
- [49] M. Tsetlin, "On behaviour of finite automata in random medium," *Avtomat. I Telemekh.*, vol. 22, no. 10, pp. 1345–1354, 1961.
- [50] M. Thathachar and P. Sastry, *Adaptive Stochastic Algorithms for Pattern Classification*. Singapore: World Sci., 2001.
- [51] T. P. I. Ahamed, P. S. N. Rao, and P. S. Sastry, "A reinforcement learning approach to automatic generation control," *Elect. Power Syst. Res.*, vol. 63, no. 1, pp. 9–26, 2002.
- [52] K. S. Narendra and M. A. L. Thathachar, *Learning Automata: An Introduction*. Courier Corporat., 2012.
- [53] B. J. Oommen and E. V. de St. Croix, "Graph partitioning using learning automata," *IEEE Trans. Comput.*, vol. 45, no. 2, pp. 195–208, Feb. 1996.
- [54] P. Du Bois *et al.*, "Vehicle crashworthiness and occupant protection," Amer. Iron Steel Inst., Southfield, MI, USA, Rep., 2004.
- [55] X. Liao, Q. Li, X. Yang, W. Zhang, and W. Li, "Multiobjective optimization for crash safety design of vehicles using stepwise regression model," *Struct. Multidiscipl. Optim.*, vol. 35, no. 6, pp. 561–569, 2008.
- [56] R. J. Yang, N. Wang, C. H. Tho, J. P. Bobineau, and B. P. Wang, "Metamodeling development for vehicle frontal impact simulation," *J. Mech. Design*, vol. 127, no. 5, pp. 1014–1020, 2005.
- [57] A. Kheiri and E. Keedwell, "A sequence-based selection hyper-heuristic utilising a hidden Markov model," in *Proc. Annu. Conf. Genet. Evol. Comput.*, Madrid, Spain, 2015, pp. 417–424.
- [58] P. Cowling, G. Kendall, and E. Soubeiga, "A hyperheuristic approach to scheduling a sales summit," in *Practice and Theory of Automated Timetabling III*. Heidelberg, Germany: Springer, 2000, pp. 176–190.
- [59] J. H. Drake, E. Özcan, and E. K. Burke, "A modified choice function hyper-heuristic controlling unary and binary operators," in *Proc. IEEE Congr. Evol. Comput. (CEC)*, Sendai, Japan, 2015, pp. 3389–3396.
- [60] J. Knowles, L. Thiele, and E. Zitzler, "A tutorial on the performance assessment of stochastic multiobjective optimizers," *ETH Zurich, Zürich, Switzerland, Tik Rep.* 214, pp. 327–332, 2006.
- [61] K. Bringmann and T. Friedrich, "The maximum hypervolume set yields near-optimal approximation," in *Proc. 12th Annu. Conf. Genet. Evol. Comput.*, Portland, OR, USA, 2010, pp. 511–518.
- [62] K. Bringmann and T. Friedrich, "Tight bounds for the approximation ratio of the hypervolume indicator," in *Proc. Int. Conf. Parallel Problem Solving Nat.*, Kraków, Poland, 2010, pp. 607–616.
- [63] R. S. Sutton and A. G. Barto, *Reinforcement Learning: An Introduction*, vol. 1. Cambridge, MA, USA: MIT Press, 1998.
- [64] J. J. Durillo and A. J. Nebro, "jMetal: A java framework for multi-objective optimization," *Adv. Eng. Softw.*, vol. 42, no. 10, pp. 760–771, 2011.
- [65] D. B. Gümüş, E. Özcan, and J. Atkin, "An investigation of tuning a memetic algorithm for cross-domain search," in *Proc. IEEE Congr. Evol. Comput. (CEC)*, Vancouver, BC, Canada, 2016, pp. 135–142.
- [66] N. S. Jaddi, S. Abdullah, and A. R. Hamdan, "Taguchi-based parameter designing of genetic algorithm for artificial neural network training," in *Proc. Int. Conf. Informat. Creative Multimedia (ICICM)*, Kuala Lumpur, Malaysia, 2013, pp. 278–281.
- [67] R. K. Roy, *A Primer on the Taguchi Method*. Dearborn, MI, USA: Soc. Manuf. Eng., 2010.
- [68] M. Wagner, K. Bringmann, T. Friedrich, and F. Neumann, "Efficient optimization of many objectives by approximation-guided evolution," *Eur. J. Oper. Res.*, vol. 243, no. 2, pp. 465–479, 2015.
- [69] T. Tušar and B. Filipič, "Differential evolution versus genetic algorithms in multiobjective optimization," in *Proc. Int. Conf. Evol. Multi Criterion Optim.*, Matsushima, Japan, 2007, pp. 257–271.

- [70] K. Deb and R. B. Agrawal, "Simulated binary crossover for continuous search space," *Complex Syst.*, vol. 9, no. 3, pp. 1–15, 1994.
- [71] B. L. Miller and D. E. Goldberg, "Genetic algorithms, tournament selection, and the effects of noise," *Complex Syst.*, vol. 9, no. 3, pp. 193–212, 1995.
- [72] L.-Y. Deng, "Orthogonal arrays: Theory and applications," *Technometrics*, vol. 42, no. 4, pp. 440–440, 2000.
- [73] N. P. Suh, *The Principles of Design*. New York, NY, USA: Oxford Univ. Press, 1990.
- [74] W. Li *et al.*, "A modified indicator-based evolutionary algorithm (mIBEA)," in *Proc. IEEE Congr. Evol. Comput. (CEC)*, 2017, pp. 1047–1054.
- [75] S. Bechikh, M. Elarbi, and L. B. Said, "Many-objective optimization using evolutionary algorithms: A survey," in *Recent Advances in Evolutionary Multi-Objective Optimization*. Cham, Switzerland: Springer, 2017, pp. 105–137.
- [76] K. Deb and H. Jain, "An evolutionary many-objective optimization algorithm using reference-point-based nondominated sorting approach, part I: Solving problems with box constraints," *IEEE Trans. Evol. Comput.*, vol. 18, no. 4, pp. 577–601, Aug. 2014.



Wenwen Li received the B.Eng. degree in software engineering from the Computer Science School, Anhui University, Hefei, China, in 2008, and the M.Sc. degree (with distinction) in operational research from the Mathematics School, University of Edinburgh, Edinburgh, U.K., in 2013. She is currently pursuing the Ph.D. degree with the University of Nottingham, Nottingham, U.K.

Her current research interests include multiobjective evolutionary algorithms, hyper-heuristics, and the application of machine learning methods in the design of evolutionary algorithms.



Ender Özcan (SM'14) received the Ph.D. degree from the Department of Computer and Information Science, Syracuse University, Syracuse, NY, USA, in 1998.

He is a Lecturer with the Automated Scheduling, Optimisation and Planning Research Group, School of Computer Science, University of Nottingham, Nottingham, U.K. Before being appointed to the EPSRC funded LANCS initiative in 2008, he was with Yeditepe University, Istanbul, Turkey, from 1998 to 2008, where he served as the Deputy Head of the Computer Engineering Department from 2004 to 2007. He has over 120 publications in his field of expertise. His current research interests include interface of computer science, artificial intelligence, and operational research with a focus on intelligent decision support systems combining data science techniques and (hyper/meta)heuristics applied to real world problems.

Dr. Özcan is the Co-Founder and the Co-Chair of the EURO Working Group on Data Science Meets Optimization. He is the Deputy Director of EPSRCs—A National Taught Course Centre in Operational Research (NATCOR). He is an Associate Editor of the *Journal of Scheduling*, the *IEEE TRANSACTIONS ON EMERGING TOPICS IN COMPUTATIONAL INTELLIGENCE*, and the *Journal of Applied Metaheuristic Computing*. He is a member of the EPSRC Peer Review College. For more details, visit <http://www.cs.nott.ac.uk/pszeo/>.



Robert John (SM'10) received the Ph.D. degree in type-2 fuzzy systems from De Montfort University, Leicester, U.K., in 2000.

He joined the University of Nottingham, Nottingham, U.K., as the Head of the Automated Scheduling, Optimisation and Planning Group, School of Computer Science, in 2013. He has over 150 publications with an H-index of 32, and over 5000 citations with papers in the top 1% most cited in ISI.

Prof. John was a recipient of grants worth circa 1 million funded by the U.K. government and Horizon 2020 and grants over 3 million. He is a member of the EPSRC Peer Review College. A leading researcher in type-2 fuzzy logic. He was the Co-General Chair of the 2007 FUZZ-IEEE International Conference. He is a member of the Editorial Board of the *International Journal of Cognitive Neurodynamics*, and an Associate Editor of the *International Journal of Information and Systems Sciences and Soft Computing*. For more details, visit <http://www.cs.nott.ac.uk/rjij>.